

Project Overview, Problem Statement & Competitive Novelty

Project: Deep-Sea Rescue Support: Human & Object Detection Using Sonar and AI Techniques

1. Formal Problem Statement

Deep-sea maritime emergencies—such as downed aircraft, sunken vessels, submerged vehicles, or missing divers—present an environment of **extreme depth, severe time urgency, and near-zero visibility**.

1.1 The Optical Failure at Depth

- Water absorbs and scatters optical light exponentially. Below depths of 30–50 meters, or in turbid coastal and post-disaster waters, optical underwater cameras degrade to an effective visibility range of **less than 1 to 2 meters**.
- Optical cameras require heavy artificial lighting which causes severe backscatter, reflecting off suspended sediment and blinding the camera sensors.

1.2 The Acoustic Dilemma

- Underwater acoustic sensors—specifically **Multibeam Forward-Looking Sonar (MFLS)**—solve the visibility problem by emitting acoustic pulses (750 kHz to 1.2 MHz) that travel through muddy water, providing a wide fan-shaped field of view up to 50–100 meters.
- However, acoustic sonar imagery is **fundamentally difficult for humans to interpret**:
 1. Low signal-to-noise ratio (SNR) with heavy speckle noise.
 2. Non-uniform acoustic shadows and multipath reverberation artifacts.
 3. Targets lack optical features (no color or fine texture); shapes appear as distorted intensity highlights with dark acoustic shadows.

1.3 The Operator Cognitive Bottleneck (The Critical Pain Point)

- In a life-or-death rescue mission, human divers have a survivability window dictated by oxygen reserves and hypothermia (often under 60–120 minutes).
- Currently, shipborne rescue teams and Remotely Operated Vehicle (ROV) pilots must **manually stare at raw, noisy, grayscale sonar sweeps** for hours.
- Operators must simultaneously:
 1. Scan for faint target shapes (human bodies, wreckage, liferafts).

2. Estimate how far away the target is and in which direction the ROV must steer.
 3. Decide which target among dozens of clutter items (rocks, tires, sea cucumbers) to investigate first.
- This creates **extreme cognitive overload**, leading to delayed interventions, missed detections, and tragic loss of life.

1.4 The Academic Research Gap

Existing academic literature suffers from severe fragmentation:

1. **Isolated Detection Studies:** Computer vision papers focus solely on improving mean Average Precision (mAP) on static datasets, outputting raw bounding boxes without giving rescue operators actionable guidance.
2. **Disconnected Localization:** Robotics and SLAM papers study acoustic range/bearing estimation for obstacle avoidance, but do not combine this with deep learning detectors to identify *what* the obstacle is.
3. **Zero Mission Prioritization:** No prior published work systematically calculates a **mission-oriented priority score** to rank which detection must be rescued first.

2. How This Project is Different from ALL Existing Work (The Novelty & USP)

Our system is **not just another object detector**. It is an **integrated perception-to-action decision support framework**.

Below is a detailed comparison of our system against prior research and commercial sonar software:

Dimension	Classical Sonar Software (Oculus View / Tritech Genesis)	SOTA Academic Models (Yu et al., STAFNet, CCW-YOLOv5)	Our Proposed System (SRM Minor Project)
Core Function	Visualizes raw acoustic echo sweeps	Classifies & bounds objects on static benchmark frames	Full-Stack Perception, Metric Localization & Rescue Triage
Target Identification	Manual (Human eyes only)	Automated via CNN / Transformers	Automated via Optimized YOLOv8 (Anchor-free, C2f modules)
Spatial Localization	Manual cursor distance measurement	None (outputs pixel coordinates only: \$x_1, y_1, x_2, y_2\$)	Automated Geometry-Aware Metric Mapping (Distance in meters, Bearing in degrees)
Target Triage / Ranking	None	None (all bounding boxes have equal importance)	Multi-Factor Priority Scoring (Eq. 3) ranking targets by rescue urgency
Operator Interface	Complex raw engineering display	None (command-line evaluation script)	Next.js Real-Time Mission Dashboard + Standalone Tkinter HUD
Temporal Tracking	None	None	IoU Multi-Object Tracker assigning persistent Track IDs & velocity
Uncertainty Estimation	None	None	Monte-Carlo Test-Time Augmentation (MC-TTA) confidence std. dev.
Mission Persistence	Manual screen capture	None	SQLite Database storing full mission sessions & target logs

3. The 4 Pillars of Novelty (Our Unique Contributions)

Pillar 1: Coupled AI Detection with Geometry-Aware Localization

Previous deep learning models output bounding box coordinates in pixels (\$u, v\$). A rescue operator steering a vehicle cannot use "pixel (412, 180)".

Our system couples the detector with the physical geometry of multibeam forward-looking sonar, projecting pixel centroids \$(u_i, v_i)\$ into real-world spherical coordinates relative to the vehicle's acoustic center:

$$\hat{r}_i = \left(1 - \frac{v_i}{H}\right) \times r_{\text{max}} + r_{\text{blind}}$$

$$\hat{\theta}_i = \left(\frac{u_i}{W} - 0.5 \right) \times \theta_{\text{FoV}}$$

- **Where:** H , W are frame dimensions, $r_{\text{max}} = 10\text{m}$, $r_{\text{blind}} = 1\text{m}$, and $\theta_{\text{FoV}} = 130^\circ$.
- **Impact:** Every detected human, ROV, or debris item instantly displays: `"Range: 4.8m | Bearing: +14.2°"`.

Pillar 2: Multi-Factor Rescue Priority Scoring (Equation 3)

In a catastrophic maritime disaster, an image may contain multiple targets: submerged wreckage, tires, steel drums, and a human diver. Operators cannot afford to investigate debris first.

We formulated a novel **Multi-Factor Priority Scoring Algorithm (\$P_i\$)**:

$$P_i = \left(w_{\text{class}} \cdot C_{\text{class}} + w_{\text{dist}} \cdot \left(\frac{r_{\text{max}} - \hat{r}}{r_{\text{max}}} \right) + w_{\text{angle}} \cdot \left(1 - \frac{|\hat{\theta}_i|}{\theta_{\text{FoV}} / 2} \right) \right) \times 10$$

Where:

1. **\$C_{\text{class}}\$ (Class Criticality Factor):** Assigns mission weight to object types (`Human Body: 100` , `ROV: 90` , `Plane: 80` , down to `Metal Bucket: 5` , `Tyre: 5`).
2. **Proximity Factor:** Rewards targets closer to the rescue vehicle ($1 - \hat{r}/r_{\text{max}}$) for immediate extraction.
3. **Heading Alignment Factor:** Rewards targets directly ahead along the vehicle's heading vector ($|\hat{\theta}_i| \rightarrow 0$).

Ablation Study Validation:

We validated our formula across 5 distinct weight configurations in `ablation_study.py`. The **Balanced configuration** ($w_{\text{class}}=0.50$, $w_{\text{dist}}=0.30$, $w_{\text{angle}}=0.20$) proved optimal, ensuring **100% Recall of Human Body at Rank #1** across all test frames.

Pillar 3: Operator-Centric Decision Interface (Not Just Code)

Instead of requiring operators to inspect Python terminal logs, we built a modern **Rescue Operator Console** in Next.js + React + Tailwind CSS:

- **Color-Coded Urgency Triage:**
- ● **Critical Priority (\$P_i > 7.0\$):** Flashing rescue alert (e.g., Human Body within 5m).
- ● **Medium Priority (\$4.0 \leq P_i \leq 7.0\$):** Navigational cue / operational target (e.g., ROV, sunken aircraft).
- ● **Low Priority (\$P_i < 4.0\$):** Passive marine debris (e.g., tires, cages).

- **Dynamic Weight Adjuster:** Operators can shift sliders in real time from "Human-First Mode" to "Proximity-First Mode".
 - **Session Replay & Logging:** Automatic SQLite mission persistence storing timestamps, coordinates, and candidate histories.
-

Pillar 4: Rescue-Reliable AI Optimization

Sonar imagery possesses fundamentally different acoustic properties than standard COCO photographic images:

- **Domain-Specific Augmentations:** We deliberately excluded standard image augmentations like heavy vertical flips and CutMix that break the physical acoustic shadow orientation of sonar beam propagation.
 - **Regularization & Stability:** Integrated **Label Smoothing (0.03)** and **Dropout (0.15)** in `train_optimized.py` to prevent neural overconfidence on noisy speckle patterns.
 - **Monte-Carlo Uncertainty:** Integrated 2-pass Test-Time Augmentation (MC-TTA) to return confidence standard deviation (σ), alerting the rescue pilot when the AI is uncertain about a degraded contact.
-

4. End-to-End System Architecture Pipeline

The system operates across **5 sequential stages**:

[Raw Sonar Frame (BMP/Sensor Stream)]



STAGE 1: Dataset & Acquisition Ingestion

- UATD Benchmark: 10 object classes
- Custom XML-to-YOLO normalized conversion



STAGE 2: Acoustic Image Preprocessing & Augmentation

- Contrast normalization, speckle noise reduction
- Scale, translation, mosaic, and HSV color jitter



STAGE 3: YOLOv8 Object Detection

- Modified CSPDarknet53 + C2f gradient feature fusion
- PANet + FPN bidirectional multiscale feature neck
- Decoupled anchor-free classification & box heads
- Outputs: Class ID, Confidence c_i , Bounding Box (u,v)



STAGE 4: Geometry-Aware Polar Metric Localization

- Projects image centroid (u, v) into physical space
- Range: r_i (meters) | Bearing: θ_i (degrees)



STAGE 5: Multi-Factor Priority Triage & Visualization

- Priority Scoring: $P_i = f(\text{Class}, \text{Range}, \text{Bearing})$
- IoU Temporal Tracker (persistent TRK IDs & velocity)
- Surface Next.js Dashboard & WebSocket Telemetry

5. Summary of Results

Parameter	Baseline Experiment	Optimized System	Improvement / Significance
Model Variant	YOLOv8n (Nano)	YOLOv8l (Large)	Richer feature representation
Parameter Count	3.2 Million	43.7 Million	Higher capacity for acoustic clutter
Detection Accuracy (\$mAP@0.5\$)	0.828 (82.8%)	0.88 – 0.92 (88-92%)	Significant accuracy gain
High-Precision (\$mAP@0.5:0.95\$)	0.361 (36.1%)	0.45 – 0.55 (45-55%)	Sharper bounding boxes
Human Recall-at-Top-1	Not calculated	100%	Lifesaving target prioritized first
Inference Latency	12 ms	28 ms (TensorRT)	Real-time capable (>30 FPS)

6. Viva / Project Defense FAQ (Questions Professors Will Ask)

Q1: "Why did you use YOLOv8 instead of Faster R-CNN or Swin Transformers?"

Answer: "Two-stage detectors like Faster R-CNN have inference latencies exceeding 80–120ms, which is too slow for real-time robotic collision avoidance and rescue tracking. While Transformers (like STAFNet) show high accuracy, their self-attention computational complexity ($O(N^2)$) exceeds the thermal and power envelope of subsea edge hardware like the NVIDIA Jetson Orin (15W). YOLOv8's anchor-free decoupled head and C2f feature extraction offer the ideal Pareto-optimal balance: high small-target sensitivity at real-time speeds (>30 FPS)."

Q2: "How can you estimate distance and angle without depth or 3D lidar?"

Answer: "Multibeam forward-looking sonar operates on known polar acoustic imaging geometry. The vertical axis directly maps to acoustic time-of-flight (range from 1m blind zone to 10m maximum range), while the horizontal axis corresponds to the transducer array's angular beam aperture (-65° to $+65^\circ$). By mapping pixel coordinates to these physical acoustic parameters, we achieve deterministic metric localization without requiring expensive subsea stereo cameras."

Q3: "What happens if a rock or debris looks identical to a human body?"

Answer: "We address acoustic ambiguity in three ways: First, our training pipeline applies domain-specific augmentations (mosaic, mixup, HSV jitter) to capture high-order shadow and edge gradients. Second, we integrate Monte-Carlo Test-Time Augmentation (MC-TTA) to calculate confidence uncertainty (σ). If uncertainty is high, the system flags the contact as 'Unverified' rather than giving false certainty. Third, our IoU temporal tracker validates candidate consistency across sequential frames."

Q4: "What is your contribution compared to the UATD dataset authors?"

Answer: "The UATD authors (Ye et al., 2022) created an annotated dataset for benchmark detection. Our contribution is the complete operational bridge: we took raw detections, developed the geometry-aware localization algorithm, formulated the multi-factor rescue priority scoring system, proved it through ablation testing, and deployed it inside a production-grade FastAPI backend and real-time Next.js operator console for deep-sea search and rescue teams."